

## Assignment 9.1

### Advance Hive

#### Task 1

1. Write a Hive program to find the number of medals won by each country in swimming.

We have created **olympics** table with all required fields as columns and as we are using csv file '**olympix\_data**' to load data into this table, we are using ',' as delimiter to separate column values:

```
hive> create table olympics
> (athlete string,
> age tinyint,
> country string,
> year smallint,
> closing_date string,
> sport string,
> gold_medals smallint,
> silver_medals smallint,
> bronze_medals smallint,
> total_medals smallint)
> row format delimited fields terminated by ',';
OK
Time taken: 3.176 seconds
hive> select * from olympics;
```

Then we have loaded data from '**olympix\_data.csv**' file present in local system as shown in below screenshot.

```
hive> load data local inpath '/home/acadgild/olympix_data.csv' into table olympics;
Loading data to table custom.olympics
OK
Time taken: 1.217 seconds
```

```
hive> set hive.cli.print.header = true;
```

#### Hive Query 1:

Here we are using **GROUP BY** clause to group all the records by using column 'country' and **sum** function to calculate total number of medals for a country in Swimming sport.

## Assignment 9.1

### Advance Hive

select sum(total\_medals) Total\_Medals, country from olympics where sport='Swimming' group by country;

```
hive> select sum(total_medals) Total_Medals, country from olympics where sport='Swimming' group by country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180811131343_a2b78b2a-212a-468d-9823-953ddd6624d8
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1533965770339_0012, Tracking URL = http://localhost:8088/proxy/application_1533965770339_0012/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1533965770339_0012
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-08-11 13:14:01,241 Stage-1 map = 0%, reduce = 0%
2018-08-11 13:14:17,937 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.53 sec
2018-08-11 13:14:38,538 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.53 sec
MapReduce Total cumulative CPU time: 8 seconds 70 msec
Ended Job = job_1533965770339_0012
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 8.07 sec HDFS Read: 536679 HDFS Write: 881 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 70 msec
OK
```

Output of Query 1:

```
OK
total_medals    country
1              Argentina
163            Australia
3              Austria
2              Belarus
8              Brazil
5              Canada
35             China
2              Costa Rica
1              Croatia
1              Denmark
39             France
32             Germany
11             Great Britain
9              Hungary
16             Italy
43             Japan
1              Lithuania
46             Netherlands
2              Norway
3              Poland
6              Romania
20             Russia
1              Serbia
2              Slovakia
1              Slovenia
11             South Africa
4              South Korea
3              Spain
9              Sweden
1              Trinidad and Tobago
3              Tunisia
7              Ukraine
267            United States
7              Zimbabwe
Time taken: 56.876 seconds, Fetched: 34 row(s)
```

## Assignment 9.1

### Advance Hive

#### 2. Write a Hive program to find the number of medals that India won year wise.

Here we are using **GROUP BY** clause to group all the records by using column 'year' and **sum** function to calculate total number of medals for a year in **India** country.

#### Hive Query 2:

**select sum(total\_medals) Total\_Medals ,year from olympics where country = 'India' group by year;**

```
hive> select sum(total_medals) Total_Medals ,year from olympics where country = 'India' group by year;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180811131924_14de2f2f-c13c-4f86-8ec6-aebc30edd185
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1533965770339_0013, Tracking URL = http://localhost:8088/proxy/application_1533965770339_0013/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1533965770339_0013
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-08-11 13:19:39,335 Stage-1 map = 0%, reduce = 0%
2018-08-11 13:19:53,306 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.03 sec
2018-08-11 13:20:09,948 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.92 sec
MapReduce Total cumulative CPU time: 7 seconds 920 msec
Ended Job = job_1533965770339_0013
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.92 sec HDFS Read: 536695 HDFS Write: 163 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 920 msec
ok
total_medals    year
1              2000
1              2004
3              2008
6              2012
Time taken: 48.117 seconds, Fetched: 4 row(s)
hive>
```

#### Output :

**total\_medals year**

|   |      |
|---|------|
| 1 | 2000 |
| 1 | 2004 |
| 3 | 2008 |
| 6 | 2012 |

## Assignment 9.1

### Advance Hive

#### 3. Write a Hive Program to find the total number of medals each country won.

Here we are using **GROUP BY** clause to group all the records by using column 'country' and **sum** function to calculate total number of medals for a country.

#### Hive Query 3:

**select sum(total\_medals) Total\_Medals ,country from olympics group by country;**

```
hive> select sum(total_medals) Total_Medals ,country from olympics group by country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180811132911_d4078c5e-2ef6-49d1-8c15-4eb7066056be
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1533965770339_0014, Tracking URL = http://localhost:8088/proxy/application_1533965770339_0014/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1533965770339_0014
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-08-11 13:29:27,546 Stage-1 map = 0%, reduce = 0%
2018-08-11 13:29:39,193 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.65 sec
2018-08-11 13:29:53,243 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.13 sec
MapReduce Total cumulative CPU time: 6 seconds 130 msec
Ended Job = job_1533965770339_0014
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.13 sec HDFS Read: 535861 HDFS Write: 2742 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 130 msec
OK
```

#### Output :

## Assignment 9.1

### Advance Hive

```
total_medals  country
2            Afghanistan
8            Algeria
141          Argentina
10           Armenia
609          Australia
91           Austria
25           Azerbaijan
24           Bahamas
1            Bahrain
1            Barbados
97           Belarus
18           Belgium
1            Botswana
221          Brazil
41           Bulgaria
20           Cameroon
370          Canada
22           Chile
530          China
20           Chinese Taipei
13           Colombia
2            Costa Rica
81           Croatia
188          Cuba
1            Cyprus
81           Czech Republic
89           Denmark
5            Dominican Republic
1            Ecuador
8            Egypt
1            Eritrea
18           Estonia
29           Ethiopia
118          Finland
318          France
```

## Assignment 9.1

### Advance Hive

|     |               |
|-----|---------------|
| 1   | Gabon         |
| 23  | Georgia       |
| 629 | Germany       |
| 322 | Great Britain |
| 59  | Greece        |
| 1   | Grenada       |
| 1   | Guatemala     |
| 3   | Hong Kong     |
| 145 | Hungary       |
| 15  | Iceland       |
| 11  | India         |
| 22  | Indonesia     |
| 24  | Iran          |
| 9   | Ireland       |
| 4   | Israel        |
| 331 | Italy         |
| 80  | Jamaica       |
| 282 | Japan         |
| 42  | Kazakhstan    |
| 39  | Kenya         |
| 2   | Kuwait        |
| 3   | Kyrgyzstan    |
| 17  | Latvia        |
| 30  | Lithuania     |
| 1   | Macedonia     |
| 3   | Malaysia      |
| 1   | Mauritius     |
| 38  | Mexico        |
| 5   | Moldova       |
| 10  | Mongolia      |
| 14  | Montenegro    |
| 11  | Morocco       |
| 1   | Mozambique    |
| 318 | Netherlands   |
| 52  | New Zealand   |
| 39  | Nigeria       |
| 21  | North Korea   |

## Assignment 9.1

### Advance Hive

|      |                       |
|------|-----------------------|
| 192  | Norway                |
| 1    | Panama                |
| 17   | Paraguay              |
| 80   | Poland                |
| 9    | Portugal              |
| 2    | Puerto Rico           |
| 3    | Qatar                 |
| 123  | Romania               |
| 768  | Russia                |
| 6    | Saudi Arabia          |
| 31   | Serbia                |
| 38   | Serbia and Montenegro |
| 7    | Singapore             |
| 35   | Slovakia              |
| 25   | Slovenia              |
| 25   | South Africa          |
| 308  | South Korea           |
| 205  | Spain                 |
| 1    | Sri Lanka             |
| 1    | Sudan                 |
| 181  | Sweden                |
| 93   | Switzerland           |
| 1    | Syria                 |
| 3    | Tajikistan            |
| 18   | Thailand              |
| 1    | Togo                  |
| 19   | Trinidad and Tobago   |
| 4    | Tunisia               |
| 28   | Turkey                |
| 1    | Uganda                |
| 143  | Ukraine               |
| 1    | United Arab Emirates  |
| 1312 | United States         |
| 1    | Uruguay               |
| 19   | Uzbekistan            |
| 4    | Venezuela             |
| 2    | Vietnam               |
| 7    | Zimbabwe              |

## Assignment 9.1

### Advance Hive

#### 4. Write a Hive program to find the number of gold medals each country won.

Here we are using **GROUP BY** clause to group all the records by using column 'country' and **sum** function to calculate total number of Gold medals for a country.

#### Hive Query 4 :

```
select sum(gold_medals) Gold_Medals ,country from olympics group by country;
```

```
hive> select sum(gold_medals) Gold_Medals ,country from olympics group by country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180811134058_957864af-002c-416a-82ac-03ebcd80175e
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1533965770339_0015, Tracking URL = http://localhost:8088/proxy/application_1533965770339_0015/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1533965770339_0015
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-08-11 13:41:15,164 Stage-1 map = 0%, reduce = 0%
2018-08-11 13:41:31,386 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.04 sec
2018-08-11 13:41:46,887 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.65 sec
MapReduce Total cumulative CPU time: 7 seconds 650 msec
Ended Job = job_1533965770339_0015
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.65 sec HDFS Read: 535852 HDFS Write: 2703 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 650 msec
OK
```



## Assignment 9.1

### Advance Hive

```
gold_medals    country
0      Afghanistan
2      Algeria
49     Argentina
0      Armenia
163    Australia
36     Austria
6      Azerbaijan
11     Bahamas
0      Bahrain
0      Barbados
17     Belarus
2      Belgium
0      Botswana
46     Brazil
8      Bulgaria
20     Cameroon
168    Canada
3      Chile
234    China
2      Chinese Taipei
2      Colombia
0      Costa Rica
35     Croatia
57     Cuba
0      Cyprus
14     Czech Republic
46     Denmark
3      Dominican Republic
0      Ecuador
1      Egypt
0      Eritrea
6      Estonia
13     Ethiopia
11     Finland
108    France
0      Gabon
```

## Assignment 9.1

### Advance Hive

|     |               |
|-----|---------------|
| 6   | Georgia       |
| 223 | Germany       |
| 124 | Great Britain |
| 12  | Greece        |
| 1   | Grenada       |
| 0   | Guatemala     |
| 0   | Hong Kong     |
| 77  | Hungary       |
| 0   | Iceland       |
| 1   | India         |
| 5   | Indonesia     |
| 10  | Iran          |
| 1   | Ireland       |
| 1   | Israel        |
| 86  | Italy         |
| 24  | Jamaica       |
| 57  | Japan         |
| 13  | Kazakhstan    |
| 11  | Kenya         |
| 0   | Kuwait        |
| 0   | Kyrgyzstan    |
| 3   | Latvia        |
| 5   | Lithuania     |
| 0   | Macedonia     |
| 0   | Malaysia      |
| 0   | Mauritius     |
| 19  | Mexico        |
| 0   | Moldova       |
| 2   | Mongolia      |
| 0   | Montenegro    |
| 2   | Morocco       |
| 1   | Mozambique    |
| 101 | Netherlands   |
| 18  | New Zealand   |
| 6   | Nigeria       |
| 6   | North Korea   |
| 97  | Norway        |

## Assignment 9.1

### Advance Hive

```
1 Panama
0 Paraguay
20 Poland
1 Portugal
0 Puerto Rico
0 Qatar
57 Romania
234 Russia
0 Saudi Arabia
1 Serbia
11 Serbia and Montenegro
0 Singapore
10 Slovakia
5 Slovenia
10 South Africa
110 South Korea
19 Spain
0 Sri Lanka
0 Sudan
57 Sweden
21 Switzerland
0 Syria
0 Tajikistan
6 Thailand
0 Togo
1 Trinidad and Tobago
2 Tunisia
9 Turkey
1 Uganda
31 Ukraine
1 United Arab Emirates
552 United States
0 Uruguay
5 Uzbekistan
1 Venezuela
0 Vietnam
2 Zimbabwe
Time taken: 49.215 seconds, Fetched: 110 row(s)
```

## Assignment 9.1

### Advance Hive

#### Task 2 :

Write a hive UDF that implements functionality of string `concat_ws(string SEP, array<string>)`. This UDF will accept two arguments, one string and one array of string. It will return a single string where all the elements of the array are separated by the SE

We have written below java program '`concatenate.java`' .

**Evaluate** function takes delimiter SEP and Array of Strings (List, in Java) as input. Then returns **word** as concatenated string like output of `concat_ws` function.

```
import org.apache.hadoop.hive.ql.exec.UDF;
import org.apache.hadoop.io.Text;
import java.util.List;

public class concatenate extends UDF {

    public Text evaluate(Text SEP,List<String> arr) {
        Text to_value = new Text("");
        if (arr != null) {
            String word = "";
            for (int i=0; i<arr.size(); i++) {
                if(i==0)
                    word =word + arr.get(i);
                else
                    word = word+ SEP + arr.get(i);
            }

            to_value.set(word);
        }
        return to_value;
    }
}
```

Then from this java code, we have exported JAR file as '`concatenate.jar`'.

After this we have added this JAR file into our VM in path location: `/home/acadgild/hive`

```
[acadgild@localhost ~]$ ls -l /home/acadgild/hive
total 32
-rw-rw-r-- 1 acadgild acadgild 29069 Aug 12 19:35 concatenate.jar
```

Then we have created a temporary function `concatenate` below :

```
hive> ADD JAR /home/acadgild/hive/concatenate.jar;
Added [/home/acadgild/hive/concatenate.jar] to class path
Added resources: [/home/acadgild/hive/concatenate.jar]
```

```
hive> CREATE TEMPORARY FUNCTION concatenate as 'concatenate';
OK
Time taken: 0.009 seconds
```

## Assignment 9.1

### Advance Hive

We have created table **array\_concat** with columns as sep as string and line as array of strings.

```
hive> create table array_concat
> (sep string,
> line array<string>)
> row format delimited
> fields terminated by ';'
> collection items terminated by ',';
OK
Time taken: 1.091 seconds
hive> select * from array_concat;
OK
Time taken: 0.448 seconds
```

Below is content in array.txt with two fields:

```
[acadgild@localhost ~]$ cat array.txt
-;Sachin,Gorade,Mumbai
*;Acadgild,online,Hadoop,course
~;This,is,Hive,Session
|;We,are,using,UDF,Function,to,replace,concat_ws
```

Then we have loaded data from array.txt file into **array\_concat** table.

```
hive> load data local inpath '/home/acadgild/array.txt' into table array_concat;
Loading data to table custom.array_concat
OK
Time taken: 2.371 seconds
hive> select * from array_concat;
OK
-      ["Sachin","Gorade","Mumbai"]
*      ["Acadgild","online","Hadoop","course"]
~      ["This","is","Hive","Session"]
|      ["We","are","using","UDF","Function","to","replace","concat_ws"]
Time taken: 0.35 seconds, Fetched: 4 row(s)
```

This is the output. You could see below that we have used **concatenate** function and used sep and line columns as input from **array\_concat** table.

e.g. in first row, '-' is delimiter and strings in array are "Sachin","Gorade" and "Mumbai".

So concatenated output is **Sachin-Gorade-Mumbai**.

```
hive> select concatenate(sep,line) from array_concat;
OK
Sachin-Gorade-Mumbai
Acadgild*online*Hadoop*course
This~is~Hive~Session
We|are|using|UDF|Function|to|replace|concat_ws
Time taken: 0.317 seconds, Fetched: 4 row(s)
```

## Assignment 9.1

### Advance Hive

#### Task 3

Link: <https://acadgild.com/blog/transactions-in-hive/>

Refer the above given link for transactions in Hive and implement the operations given in the blog using your own sample data set and send us the screenshot.

We are setting below properties in Hive. Because without setting these properties 'Update' and 'Delete' will not work and we will receive errors.

```
hive> set hive.support.concurrency = true;
hive> set hive.enforce.bucketing = true;
hive> set hive.exec.dynamic.partition.mode = nonstrict;
hive> set hive.txn.manager = org.apache.hadoop.hive.ql.lockmgr.DbTxnManager;
hive> set hive.compactor.initiator.on = true;
hive> set hive.compactor.worker.threads = 4;
hive>
```

We have created a table with name 'college' and its columns are clg\_id, clg\_name, clg\_loc. We are bucketing this table by clg\_id column and using ORC file format.

```
hive> CREATE TABLE college(clg_id int,clg_name string,clg_loc string) clustered by (clg_id) into 5 buckets stored as orc TBLPROPERTIES('transactional'='true');
OK
Time taken: 1.449 seconds
hive> show tables;
OK
array_demo
array_demo2
college
olympics
Time taken: 0.157 seconds, Fetched: 4 row(s)
hive> select * from college;
OK
Time taken: 0.482 seconds
```

Then we have inserted data into this college table with below insert command:

```
hive> INSERT INTO table college values(1,'nec','nlr'),(2,'vit','vlr'),(3,'srm','chen'),(4,'lpu','del'),(5,'stanford','uk'),(6,'JNTUA','atp'),(7,'cambridge','us');
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180812134732_6b7f69e6-6c4d-4cfb-bb79-Safaecf0b2f7
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1534050476803_0004, Tracking URL = http://localhost:8088/proxy/application_1534050476803_0004/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1534050476803_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 5
2018-08-12 13:47:50,807 Stage-1 map = 0%, reduce = 0%
2018-08-12 13:48:06,498 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.6 sec
2018-08-12 13:48:58,247 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 8.48 sec
2018-08-12 13:48:59,916 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 11.4 sec
2018-08-12 13:49:29,068 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 26.46 sec
2018-08-12 13:49:30,698 Stage-1 map = 100%, reduce = 92%, Cumulative CPU 35.66 sec
2018-08-12 13:49:34,143 Stage-1 map = 100%, reduce = 93%, Cumulative CPU 36.85 sec
2018-08-12 13:49:35,712 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 39.93 sec
MapReduce Total cumulative CPU time: 39 seconds 930 msec
Ended Job = job_1534050476803_0004
Loading data to table custom.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 5 Cumulative CPU: 39.93 sec HDFS Read: 27056 HDFS Write: 3996 SUCCESS
Total MapReduce CPU Time Spent: 39 seconds 930 msec
OK
```

## Assignment 9.1

### Advance Hive

We could see below that 7 rows have been inserted into college table successfully.

```
hive> select * from college;
OK
5      stanford      uk
6      JNTUA      atp
1      nec      nlr
7      cambridge      us
2      vit      vlr
3      srm      chen
4      lpu      del
Time taken: 0.379 seconds, Fetched: 7 row(s)
```

Now we are inserting same records again into college table and these rows will be appended.

```
hive> INSERT INTO table college values(1,'nec','nlr'),(2,'vit','vlr'),(3,'srm','chen'),(4,'lpu','del'),(5,'stanford','uk'),(6,'JNTUA','atp'),(7,'cambridge','us');
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180812135521_5c5fb945-6b58-48b3-98a4-d29449c9c803
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1534050476803_0006, Tracking URL = http://localhost:8088/proxy/application_1534050476803_0006/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1534050476803_0006
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 5
2018-08-12 13:55:37,891 Stage-1 map = 0%, reduce = 0%
2018-08-12 13:55:51,994 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.04 sec
2018-08-12 13:56:34,915 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 6.33 sec
2018-08-12 13:56:41,130 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 9.23 sec
2018-08-12 13:56:42,556 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 10.33 sec
2018-08-12 13:56:44,071 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 11.88 sec
2018-08-12 13:57:00,859 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 20.37 sec
2018-08-12 13:57:04,036 Stage-1 map = 100%, reduce = 84%, Cumulative CPU 24.76 sec
2018-08-12 13:57:05,410 Stage-1 map = 100%, reduce = 87%, Cumulative CPU 25.35 sec
2018-08-12 13:57:07,777 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 33.58 sec
MapReduce Total cumulative CPU time: 33 seconds 580 msec
Ended Job = job_1534050476803_0006
Loading data to table custom.College
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 5 Cumulative CPU: 33.58 sec HDFS Read: 26841 HDFS Write: 3995 SUCCESS
Total MapReduce CPU Time Spent: 33 seconds 580 msec
OK
```

## Assignment 9.1

### Advance Hive

We could see data in college table below :

```
hive> select * from college;
OK
5      stanford      uk
5      stanford      uk
6      JNTUA      atp
1      nec      nlr
6      JNTUA      atp
1      nec      nlr
7      cambridge      us
2      vit      vlr
7      cambridge      us
2      vit      vlr
3      srm      chen
3      srm      chen
4      lpu      del
4      lpu      del
Time taken: 0.293 seconds, Fetched: 14 row(s)
hive>
```

Below we are trying to update bucketed column 'clg\_id'. But we have received error.

So it means that we cannot update bucketed column.

```
hive> UPDATE college set clg_id = 8 where clg_id = 7;
FAILED: SemanticException [Error 10302]: Updating values of bucketing columns is not supported. Column clg_id.
```

Below we have performed update on non-bucketed column 'clg\_name' and it has been updated successfully.

This means that we can update non-bucketed column.

```
hive> UPDATE college set clg_name = 'IIT' where clg_id = 6;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180812145913_dc0775d4-b2a5-4884-bbaf-8cd79b8ba568
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1534050476803_0008, Tracking URL = http://localhost:8088/proxy/application_1534050476803_0008/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1534050476803_0008
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2018-08-12 14:59:32,577 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 5.7 sec
2018-08-12 15:00:33,297 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 22.27 sec
2018-08-12 15:00:52,737 Stage-1 map = 40%, reduce = 0%, Cumulative CPU 25.41 sec
2018-08-12 15:00:56,220 Stage-1 map = 60%, reduce = 0%, Cumulative CPU 27.51 sec
2018-08-12 15:01:59,811 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 32.35 sec
2018-08-12 15:02:03,199 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 39.16 sec
2018-08-12 15:02:16,458 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 48.71 sec
2018-08-12 15:02:17,722 Stage-1 map = 100%, reduce = 93%, Cumulative CPU 51.05 sec
MapReduce Total cumulative CPU time: 51 seconds 50 msec
Ended Job = job_1534050476803_0008
Loading data to table custom.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 51.05 sec HDFS Read: 56442 HDFS Write: 959 SUCCESS
Total MapReduce CPU Time Spent: 51 seconds 50 msec
OK
Time taken: 188.649 seconds
```



## Assignment 9.1

### Advance Hive

Below you could see that clg\_name has been changed to IIT for clg\_id =6

```
hive> select * from college;
OK
5      stanford      uk
5      stanford      uk
6      IIT      atp
1      nec      nlr
6      IIT      atp
1      nec      nlr
7      cambridge      us
2      vit      vlr
7      cambridge      us
2      vit      vlr
3      srm      chen
3      srm      chen
4      lpu      del
4      lpu      del
Time taken: 0.495 seconds, Fetched: 14 row(s)
```

Below we have deleted data having clg\_id = 1.

```
hive> delete from college where clg_id=1;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180812151007_0d2d84bd-5913-4140-b720-10c4bcf02528
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1534050476803_0009, Tracking URL = http://localhost:8088/proxy/application_1534050476803_0009/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1534050476803_0009
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2018-08-12 15:10:27,037 Stage-1 map = 0%, reduce = 0%
2018-08-12 15:11:28,256 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 9.75 sec
2018-08-12 15:11:58,111 Stage-1 map = 80%, reduce = 0%, Cumulative CPU 31.26 sec
2018-08-12 15:11:59,764 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 34.87 sec
2018-08-12 15:12:54,272 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 37.84 sec
2018-08-12 15:12:55,660 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 39.42 sec
2018-08-12 15:12:57,144 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 41.12 sec
2018-08-12 15:12:58,580 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 42.8 sec
2018-08-12 15:13:06,025 Stage-1 map = 100%, reduce = 73%, Cumulative CPU 45.12 sec
2018-08-12 15:13:07,544 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 47.07 sec
2018-08-12 15:13:08,917 Stage-1 map = 100%, reduce = 93%, Cumulative CPU 51.38 sec
2018-08-12 15:13:10,080 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 53.62 sec
MapReduce Total cumulative CPU time: 53 seconds 620 msec
Ended Job = job_1534050476803_0009
Loading data to table custom.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 53.62 sec HDFS Read: 54614 HDFS Write: 746 SUCCESS
Total MapReduce CPU Time Spent: 53 seconds 620 msec
OK
Time taken: 185.272 seconds
```

## Assignment 9.1

### Advance Hive

We could see that clg\_id having value as 1 has been deleted successfully from college table.

```
hive> select * from college;
OK
5      stanford      uk
5      stanford      uk
6      IIT      atp
6      IIT      atp
7      cambridge      us
2      vit      vlr
7      cambridge      us
2      vit      vlr
3      srm      chen
3      srm      chen
4      lpu      del
4      lpu      del
Time taken: 0.499 seconds, Fetched: 12 row(s)
```